

## Module 18 – Database Tuning (and Other Performance Issues)

### Objectives

- Identify how poor system performance results from a poor database system design.
  - Familiarize with how database tuning focuses on identifying and fixing underlying flaws.
  - Familiarize with typical objects that can be tuned.
- 

## TUNING APPLICATION DESIGN

### Effective Table Design

- Poor table design always leads to poor performance.
- Rigid adherence to **fully normalized** relational table guidelines can also result in poor physical performance. These inefficiencies result from:
  - requiring too many **joins**.
  - failure to reflect the **normal access paths** for data.
- Queries with **large numbers of columns** that come from multiple tables can cause performance to suffer because several tables must be joined.
- Design options include:
  - **Denormalizing** 1NF, 2NF and 3NF solutions.
  - Creating small **summary tables** from large, static tables that stores data in the format in which users ask for the data - this avoids joins where data are often requested and the data doesn't change very often.
  - Separating individual tables into several tables by either **vertical partitioning** and/or **horizontal partitioning**.

### Distribution of CPU Requirements

- An Oracle database that is **CPU-Bound** (limited by CPU resources) as opposed to **Wait-Bound** (waiting on disk writes of some type) is one that's effectively designed. This means that other resources are not limiting the database.
- Schedule **long-running batch query/update programs** for off-peak hours; then run them at normal priority.
- Store data in its most appropriate place in a **distributed computing** environment - this distributes computing CPU requirements from one server to another.

- Use the **Parallel Query** option to distribute processing requirements of selected SQL statements among multiple CPUs if they are available.

## Effective Application Design

- There are two principles to follow here:
  1. limit the number of times that users access the database, and
  2. coordinate the requests of users for data. This requires you to know **how** users tend to access data.
- Try to use the **same queries** to handle similar application data requirements; this will increase the likelihood that a data requirement can be resolved by data already in the **SGA**.
- Use **snapshots**, which are **non-updatable views** of part of a database that can be distributed to support typical managerial querying.
- Create stored procedures, functions, and packages and compile them to **eliminate run-time compilation**. The parsed version may exist in the **Shared SQL Pool**.
  1. The **SQL Text** for all procedures, functions, and packages can be viewed in the **TEXT** column of the **DBA\_SOURCE** view.
  2. These objects (procedural code) are stored in the **SYSTEM** tablespace so you need to allocate more space to it -- usually double its size.

---

## TUNING SQL

- Most **SQL tuning** requires the DBA to work with an application developer.
- Most improvement in database processing will come from tuning SQL.
- The key to SQL tuning is to **minimize the search path** that a database uses to find data. For the most part this requires the creation of appropriate **indexes** that the database engine will use to find data.

## Indexes

- An **Index** enables Oracle to locate a row according to its physical location in a datafile by going to the correct file, finding the correct block, then finding the correct row within the block. Taken together, these values are like **relative record numbers** in a relative file.
- The **File ID** portion of **ROWID** can be compared to the **FILE\_ID** column of the **DBA\_DATA\_FILES** view to determine the file to which an index belongs.
- A query with no **WHERE** clause normally results in a full table scan, reading every block in a table.
- A query for specific rows may cause use of an **index**.

- The index maps logical values in a table (key columns) to their **ROWIDs** which enables location of the rows directly by their physical location.
- You may index several columns together to form a **concatenated index**.
  - A concatenated index is only used if its **leading column** is used in the query's **WHERE** clause.
  - Consider the following example:  

```
CREATE INDEX City_state_zip_ndx
ON Employee (City, State, Zip)
TABLESPACE Indexes;

SELECT *
FROM Employee
WHERE State = 'NJ';
```
- The index is **NOT** used because the **WHERE** clause value of **State** does not match the leading column (**City**) of the index.
- This example would only add overhead because the index would be maintained, even if it's not used.
- The index should be recreated with proper ordering of the component fields in the index if the query is executed often.

## Ordering Data

- While **row ordering** does not matter in relational theory, it is important to order rows as much as possible when tables are initially created, e.g. when you are porting a system into Oracle from another platform or DBMS.
- Ordered rows may enable Oracle to find needed rows while minimizing the data blocks that are retrieved where users execute queries that specify **ranges of values** (recall the **BETWEEN** operator in SQL for range of value queries).
- Consider the following example which will require fewer data blocks to be read if the records are physically ordered on the **Empno** field.

```
SELECT *
FROM Employee
WHERE Empno BETWEEN 1 and 100;
```

- You can physically **sort** table rows by SELECTing them to another file with use of the ORDER BY clause, then truncating the original table and loading the rows back into the original table.

## Clusters

- We covered indexed clusters in an earlier module.

- Another type of cluster, the **hash cluster**, stores rows in a specific location based on its value in the cluster key column.
  - Every time a row is inserted, its **cluster key value** is used to determine which block to store the row in.
  - This enables hashing directly to data blocks without use of an index.
  - The hash cluster is only used with **equivalence** queries - where the exact value stored in a column is to be found.

## Explain Plan

- The **EXPLAIN PLAN** command shows the execution path for a query and stores this information to a table (**PLAN\_TABLE**) in the database. You can then query the table. Example:

```
EXPLAIN PLAN
SET Statement_id = 'TEST'
FOR
    SELECT *
    FROM Employee
    WHERE city > 'Y%';
```

- The query above is not actually executed; rather the plan for execution is stored to the **PLAN\_TABLE**.
- Your account must have a **PLAN\_TABLE** in your schema. The script to create this table is **UTLXPLAN.SQL** and is located in the **\$ORACLE\_HOME/rdbms/admin** subdirectory.
- Query the table to produce the output that shows the execution path.

```
SELECT LPAD(' ',2*LEVEL) || operation
       || ' ' || options ||
       ' ' || object_name Path_Plan
FROM Plan_Table
WHERE Statement_id = 'TEST'
CONNECT BY PRIOR Id = Parent_id
       AND Statement_id = 'TEST'
       START WITH Id=1;
```

### Path\_Plan

```
-----
TABLE ACCESS BY ROWID EMPLOYEE
INDEX RANGE SCAN CITY_ST_ZIP_NDX
```

- The output shows that data will be accessed by ROWID through an index range scan of the named index.
- Alternatively, you can also use the **SET AUTOTRACE ON** command in **SQL\*Plus** to generate the explain plan output and trace information for every query that you run.
- Evaluate the output by ensuring that the most selective (most nearly unique) indexes are used by a query.

## TUNING MEMORY USAGE

You can use the Oracle Enterprise Manager software to analyze usage of memory by Oracle's various memory caches.

- The dictionary cache in memory is not directly sized or tuned as it is part of the **Shared SQL Pool**.
- These memory areas are managed by the **LRU** (least recently used) algorithm. You set the **Shared SQL Pool size** with the **SHARED\_POOL\_SIZE** parameter.
- If your **Shared SQL Pool** is too large, then you are wasting memory.
- The **Hit Ratio** measures how well the data buffer cache handles requests.

$$\text{Hit Ratio} = (\text{Logical Reads} - \text{Physical Reads}) / \text{Logical Reads}$$

- A **perfect ratio is 1.00** - all reads are logical reads; of course, this is generally impossible to obtain since it indicates that all the data that a system user will ever need to access is stored in the **SGA**.
- On-line transaction processing applications should have Hit Ratios in excess of **0.90**.
- If processing for the **Hit Ratio** is within tolerance, you need to check to see if you can reduce the size of the Shared SQL Pool and still maintain a good Hit Ratio.
- Add the following to the **INIT.ORA** file.

```
DB_BLOCK_LRU_STATISTICS = TRUE
```

- Shutdown the database and restart it.
- The system dictionary table **SYS.X\$KCBCH** maintains memory statistics. One row is maintained for each buffer in the buffer cache. You can query this information to determine how many buffers are not being used.
- Use the following query to determine how many cache hits (the **COUNT** column) would be lost if you reduced the number of buffers (the **INDX** column).

```
SELECT Sum(Count) Lost_Hits
FROM Sys.X$Kcbcbh
WHERE indx >= New_Number_Of_Buffers;
```

(NOTE: You supply the value in the WHERE clause)

- If you have **lost hits**, the system will require additional physical reads - the Hit Ratio for this new number of data buffers is:

```
Hit Ratio =
(Logical Reads - Physical Reads - Lost Hits)
/ Logical Reads
```

- Since running the database in a statistics gathering mode will slow it down due to the additional overhead, you should comment out the **DB\_BLOCK\_LRU\_STATISTICS** parameter after you have finished tuning and restart the database.

## TUNING DATA STORAGE

### Defragmentation of Segments

**Fragmented** tables with **multiple extents** will slow down query processing. This can also slow down the storage of new records because the database may have to dynamically combine free extents to create a new extent large enough to meet the storage parameters of the object where data are being stored.

- We know that a **segment** is created to hold data associated with a new object (index or table) when an object is created.
- The space allocated is used unless the **segment** is released (dropped) or truncated (tables only).
- It would be best if each segment was composed of a single large **INITIAL** extent - compute the size for the initial extent such that it is large enough to handle all of a segment's data.
- Use the **DBA\_SEGMENTS** data dictionary view to determine which segments are comprised of ten or more extents.

```
SELECT Tablespace_name TSName, Owner,
       Segment_Name SName,
       Segment_type SType, Extents,
```

```

        Blocks, Bytes
FROM Sys.DBA_Segments;

```

TSNAME	OWNER	SNAME	SNTYPE	EXTENTS	BLOCKS	BYTES
DATA	DBOCK	LONGTIME	TABLE	1	15	61440
DATA	DBOCK	MAGAZINE	TABLE	1	15	61440
DATA	DBOCK	MATH	TABLE	1	15	61440
DATA	DBOCK	WORKERANDSKILL	CLUSTER	2	30	122880

- To see the size of the individual extents for a segment, query the **DBA\_EXTENTS** view. Supply the type of segment you desire (TABLE, INDEX, CLUSTER, ROLLBACK, TEMPORARY, etc.).

```

SELECT Tablespace_name TSNAME, Owner,
       Segment_Name SNAME,
       Segment_Type SNTYPE, Extent_id EID,
       File_id FID, Block_id BID, Bytes, Blocks
FROM Sys.DBA_Extents
WHERE Segment_type = 'segment_name'
ORDER BY Extent_id;

```

TSNAME	OWNER	SNAME	SNTYPE	EID	FID	BID	Bytes	Blocks
DATA	DBOCK	SYS_C00890	INDEX	0	4	137	61440	15
DATA	DBOCK	SYS_C00891	INDEX	0	4	152	61440	15

- If a segment is fragmented, you can rebuild the object into a single segment by using the proper size for the storage parameters. Export the data for the segment, recreate the object, then import the data into the **INITIAL** extent.

## Defragmentation of Free Extents

A **Free Extent** is a collection of contiguous free blocks in a tablespace that are unused.

- If a segment is dropped, its extents are **deallocated** and become free, but these extents are **not** recombined with neighboring free extents.
- SMON** periodically coalesces neighboring free extents only if the default **PCTINCREASE** for a tablespace is **non-zero**.
- The **ALTER TABLESPACE *tablespace\_name* COALESCE** command can be used to force the combining of free extents.
- Your readings will list a number of scripts available in Oracle to test whether or not free space needs to be coalesced.



## Identifying chained Rows

- **Chained rows** occur when a row is updated and will no longer fit into a single data block.
- If you store rows that are **larger** than the Oracle block size, then you will cause chaining.
- Chaining affects performance because of the need for Oracle to search multiple blocks for a logical row.
- The **ANALYZE** command can be used to determine if chaining is occurring.

```
ANALYZE TABLE Table_Name  
LIST CHAINED ROWS INTO Chained_Rows;
```

- The output is stored to the **CHAINED\_ROWS** table in your schema. The **CHAINED\_ROWS** table needs to first be created by executing the **UTLCHAIN.SQL** script in the **\$ORACLE\_HOME/rdbms/admin** directory.
- If chaining is prevalent (all chained rows are listed), then rebuild the table with a higher **PCTFREE** parameter.

## Increasing the Oracle Block Size

- Oracle support different block sizes, but the most common block sizes used are **4K** and **8K**.
- Installation routines default to different block size values. For example, with our version of Oracle in a LINUX environment the system **defaults to 8K** if you do not specify the block size when you create the database.
  - Using the next higher block size value may **improve** performance of query-intensive operations by up to **50 percent**.
  - **Problem**: You must rebuild the entire database to increase the block size.
  - Improvement comes because the block header does not increase significantly leaving more space in a block for transaction information and for data rows.

## Bulk Deletes: The TRUNCATE Command

- Deleting all the rows in a table will not save any space because the segment for the table is still allocated all of the extents beyond those the first one that was allocated by the **INITIAL** parameter.
- Deleting all rows can also result in UNDO errors because the bulk delete causes a very large transaction – can lead to overwrites and the **Snapshot Too Old** error message..



- The **TRUNCATE** command resolves both problems, but you need to realize that this is a **DDL** command, not a **DML** command, so it cannot be **rolled back**.
- The **TRUNCATE** command is the fastest way to delete large volumes of data.

```
TRUNCATE TABLE Employee DROP STORAGE;
```

```
TRUNCATE CLUSTER Emp_Dept REUSE STORAGE;
```

- The above command deletes all **Employee** table rows and the **DROP STORAGE** clause de-allocates the non-INITIAL extents. The second example command is for clusters.

---

**END OF NOTES**